

CSC345/M45: Big Data & Machine Learning (support vector machine)

Prof. Xianghua Xie

x.xie@swansea.ac.uk

<http://csvision.swan.ac.uk>


224 Computational Foundry, Bay Campus

Clustering vs. Classification


- Clustering: unsupervised learning
 - Class labels of the data are unknown
 - Given data, the task is to establish the existence of classes or clusters in the data
- Classification: supervised learning
 - Supervision: data (observations) are labelled with pre-defined classes
 - The input data (training set) consists of multiple records, each of which has multiple attributes or features
 - Given training data, the task is
 - to develop an accurate description or model for each class using the features
 - and to predict categorical class label for unseen data (test data)

Classification

- Supervised learning
 - Training data (X_i, Y_i) , X_i is typically a feature vector and Y_i is the corresponding class label
 - The task of training is to find a good mapping function f
 - The derived function is then evaluated on test data (unseen data)




A classifier, a mapping, a hypothesis



From Logistic Regression to SVM

- Logistic regression cost function:

$$E(\omega) = -\frac{1}{N} \sum_{i=1}^N [y_i \log h_{\omega}(x_i) + (1 - y_i) \log(1 - h_{\omega}(x_i))] + \frac{\lambda}{2N} \omega^T \omega$$




$$h_{\omega}(x) = \frac{1}{1 + e^{-\omega^T x}}$$

From Logistic Regression to SVM

- Logistic regression cost function:

$$E(\omega) = -\frac{1}{N} \sum_{i=1}^N [y_i \log h_{\omega}(x_i) + (1 - y_i) \log(1 - h_{\omega}(x_i))] + \frac{\lambda}{2N} \omega^T \omega$$

- If $y=1$, we want $h_{\omega}(x) \approx 1$, $\omega^T x \gg 0$
- If $y=0$, we want $h_{\omega}(x) \approx 0$, $\omega^T x < 0$



$$h_{\omega}(x) = \frac{1}{1 + e^{-\omega^T x}}$$

From Logistic Regression to SVM

- Logistic regression cost function:

$$E(\omega) = -\frac{1}{N} \sum_{i=1}^N [y_i \log h_{\omega}(x_i) + (1 - y_i) \log(1 - h_{\omega}(x_i))] + \frac{\lambda}{2N} \omega^T \omega$$

↳ $\min_{\omega} -\frac{1}{N} \sum_{i=1}^N [y_i \log h_{\omega}(x_i) + (1 - y_i) \log(1 - h_{\omega}(x_i))] + \frac{\lambda}{2N} \omega^T \omega$

↳ $\min_{\omega} \frac{1}{N} \sum_{i=1}^N [y_i (-\log h_{\omega}(x_i)) + (1 - y_i)(-\log(1 - h_{\omega}(x_i)))] + \frac{\lambda}{2N} \omega^T \omega$

↳ $\min_{\omega} \sum_{i=1}^N [y_i (-\log h_{\omega}(x_i)) + (1 - y_i)(-\log(1 - h_{\omega}(x_i)))] + \frac{\lambda}{2} \omega^T \omega$

↳ $\min_{\omega} C \sum_{i=1}^N [y_i (-\log h_{\omega}(x_i)) + (1 - y_i)(-\log(1 - h_{\omega}(x_i)))] + \frac{1}{2} \omega^T \omega$

↳ $\min_{\omega} C \sum_{i=1}^N [y_i \text{cost}_1(\omega^T x_i) + (1 - y_i) \text{cost}_0(\omega^T x_i)] + \frac{1}{2} \omega^T \omega$

From Logistic Regression to SVM

- Logistic regression cost function:

$$\min_{\omega} C \sum_{i=1}^N [y_i \text{cost}_1(\omega^T x_i) + (1 - y_i) \text{cost}_0(\omega^T x_i)] + \frac{1}{2} \omega^T \omega$$

$$\begin{cases} \text{cost}_1(\omega^T x_i) = -\log h_{\omega}(x_i) & \text{if } y = 1 \\ \text{cost}_0(\omega^T x_i) = -\log(1 - h_{\omega}(x_i)) & \text{if } y = 0 \end{cases}$$


From Logistic Regression to SVM

- Logistic regression cost function:


$$\min_{\omega} C \sum_{i=1}^N [y_i \text{cost}_1(\omega^T x_i) + (1 - y_i) \text{cost}_0(\omega^T x_i)] + \frac{1}{2} \omega^T \omega$$

$$\begin{cases} \text{cost}_1(\omega^T x_i) = -\log h_{\omega}(x_i) & \text{if } y = 1 \\ \text{cost}_0(\omega^T x_i) = -\log(1 - h_{\omega}(x_i)) & \text{if } y = 0 \end{cases}$$

If $y=1$, we want $h_{\omega}(x) \approx 1$, $\omega^T x \gg 0$



If $y=0$, we want $h_{\omega}(x) \approx 0$, $\omega^T x \ll 0$



From Logistic Regression to SVM

- SVM cost function (linear case):


$$\min_{\omega} C \sum_{i=1}^N [y_i \text{cost}_1(\omega^T x_i) + (1 - y_i) \text{cost}_0(\omega^T x_i)] + \frac{1}{2} \omega^T \omega$$

From Logistic Regression to SVM


- SVM cost function (linear case):

$$\min_{\omega} C \sum_{i=1}^N [y_i \text{cost}_1(\omega^T x_i) + (1 - y_i) \text{cost}_0(\omega^T x_i)] + \frac{1}{2} \omega^T \omega$$

If $y=1$, we want $\omega^T x \geq 1$



If $y=0$, we want $\omega^T x \leq -1$




From Logistic Regression to SVM

- SVM cost function (linear case):


$$\min_{\omega} C \sum_{i=1}^N [y_i \text{cost}_1(\omega^T x_i) + (1 - y_i) \text{cost}_0(\omega^T x_i)] + \frac{1}{2} \omega^T \omega$$

$$\begin{cases} \text{cost}_1(\omega^T x_i) = \max(0, 1 - \omega^T x_i) & \text{if } y = 1 \\ \text{cost}_0(\omega^T x_i) = \max(0, 1 + \omega^T x_i) & \text{if } y = 0 \end{cases}$$

If $y=1$, we want $\omega^T x \geq 1$




If $y=0$, we want $\omega^T x \leq -1$



Enforces a larger margin!

SVM Decision Boundary: large margin classifier

linearly separable data



SVM Decision Boundary

- Recap: cost function

$$\min_{\omega} C \sum_{i=1}^N [y_i \text{cost}_1(\omega^T x_i) + (1 - y_i) \text{cost}_0(\omega^T x_i)] + \frac{1}{2} \omega^T \omega$$

- Equivalent to:

$$\min_{\omega} \frac{1}{2} \omega^T \omega = \min_{\omega} \frac{1}{2} \sum_{i=1}^N \omega_i^2$$

$$s.t. \quad \omega^T x_i \geq 1, \quad \text{if } y_i = 1$$

$$\omega^T x_i \leq -1, \quad \text{if } y_i = 0$$



SVM Decision Boundary

$$\min_{\omega} \frac{1}{2} \omega^T \omega = \min_{\omega} \frac{1}{2} \sum_{i=1}^N \omega_i^2$$

$$s. t. \quad \omega^T x_i \geq 1, \quad \text{if } y_i = 1$$


$$\omega^T x_i \leq -1, \quad \text{if } y_i = 0$$

Which one of the decision boundaries is better?



Non-linear Decision Boundary


- How to represent the decision boundary?
 - Use the linear regression technique; but note the data is nonlinear
 - Use polynomials
 - Use kernel functions



Non-linear Decision Boundary

- How to represent the decision boundary?
 - Use the linear regression technique; but note the data is nonlinear
 - ~~Use polynomials~~
 - Use kernel functions



Recap: linear regression



$$y(x, \mathbf{w}) = w_0 + w_1x + w_2x^2 + \dots + w_Mx^M = \sum_{j=0}^M w_jx^j$$

$$y(\mathbf{x}, \mathbf{w}) = \sum_{j=0}^M w_j\phi_j(\mathbf{x}) = \mathbf{w}^T\boldsymbol{\phi}(\mathbf{x})$$

$$\phi_j(x) = x^j.$$



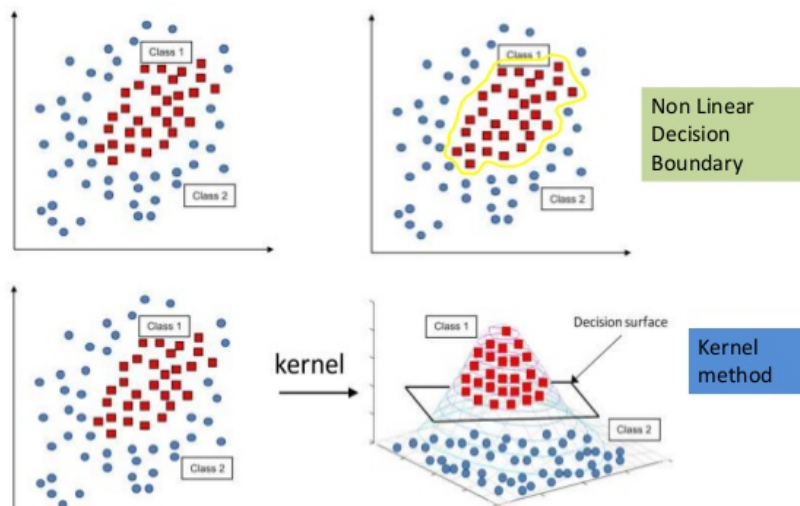
Decision Boundary: from linear to non-linear

- Linear case:

$$\min_{\omega} C \sum_{i=1}^N [y_i \text{cost}_1(\omega^T x_i) + (1 - y_i) \text{cost}_0(\omega^T x_i)] + \frac{1}{2} \omega^T \omega$$

- Non-linear case:

$$\min_{\omega} C \sum_{i=1}^N [y_i \text{cost}_1(\omega^T \phi(x_i)) + (1 - y_i) \text{cost}_0(\omega^T \phi(x_i))] + \frac{1}{2} \omega^T \omega$$



SVM Parameters

- C
 - Large C: lower bias, higher variance
 - Small C: higher bias, lower variance

$$\min_{\omega} C \sum_{i=1}^N [y_i \text{cost}_1(\omega^T \phi(x_i)) + (1 - y_i) \text{cost}_0(\omega^T \phi(x_i))] + \frac{1}{2} \omega^T \omega$$

- σ (kernel width)
 - Large σ :
 - kernel varies more smoothly
 - Higher bias, lower variance
 - small σ :
 - Kernel is sharper
 - Lower bias, higher variance

Classification

- Evaluation
 - Measure performance on independent blind test data
 - LOOCV: leave one out cross validation
 - Take one sample from the dataset as the test data and the remaining for training
 - Rotate the test data across the whole dataset
 - Performance is measured as the average across the rounds
 - K-fold cross validation:
 - Divide the dataset into K even parts
 - K-1 parts used for training, and one for testing
 - Rotating the test set
 - Performance is measured as the average across the rounds
 - 2-fold cross validation
 - Simplest variation of K-fold
 - Training and testing has equal number of samples

Classification

- Evaluation

- Positive: data sample that belongs to the class of interest
- Negative: data sample that does not belong to the class
- True positive: a positive sample correctly identified as positive
- False positive: a negative sample incorrectly classified as positive
- True negative: a negative sample correctly identified as negative
- False negative: a positive sample incorrectly classified as negative
- True positive rate: sensitivity
 - Percentage of correct prediction of positive samples $\frac{TP}{TP + FN}$
- True negative rate: specificity
 - Percentage of correct prediction of negative samples $\frac{TN}{TN + FP}$

Classification

- Evaluation
 - Overall accuracy
 - Percentage of correct prediction


$$Acc = \frac{TP + TN}{TP + FP + TN + FN} = \frac{TP + TN}{P + N}$$

- Error rate
 - Percentage of incorrect prediction
- 2-class confusion matrix

True class	Predicted class	
	positive	negative
positive	TP	?
negative	?	TN

ROC Curve


- Receiver operating characteristic (ROC), or ROC curve
 - a graphical plot that illustrates the performance of a binary classifier as its decision boundary is varied
 - X axis: false positive rate
 - Y axis: true positive rate
 - Each classifier represented by a point in ROC space corresponding to its (FP,TP) pair



- Varying parameters
- Varying class distribution


ROC Curve

- A basic ROC curve showing 5 discrete classifiers
 - Which one is the best classifier?
 - What is the dashed line?



Creating an ROC Curve

- A classifier produces a single ROC point.
- If the classifier has a “sensitivity” parameter, varying it produces a series of ROC points (confusion matrices).
- Alternatively, if the classifier is produced by a learning algorithm, a series of ROC points can be generated by varying the class ratio in the training set.
 - Class ratio (class distribution):
number of positives samples versus
number of negative samples



Example ROC Curve

- Learner L1 dominates as L2's ROC curve is beneath L1's curve
 - i.e. L1 is better than L2 for all possible costs and class distributions
- L2 & L3: neither dominates
 - Perhaps switch the classifiers at the intersection point on ROC graph

